

# SparkSocial — Feature Guide

Every feature that is built, what it is for, **the steps to test it**, and how to tell it worked.

Each entry is the same shape:

**What it is** — one line. *Story* — whose job it does. **Where** — the screen. **Steps** — numbered, do these. **Pass when** — what you should see.

Work through it in order and it doubles as a first run of the product: §1 creates a brand, and everything after it has something to act on.

## Before you start

1. API answers at `http://localhost:8080/health` with `status: ok`.
2. Web app loads at `http://localhost:3000`.
3. Migrations are applied: `npm run migrate -w @sparksocial/db`.
4. Have `docs/TEST_BUSINESSES.md` open — three businesses with every answer ready to paste.

## 1. Getting started

### 1.1 Brands and workspaces

**What it is** — one account holds many brands; each has its own content, assets, connections and rules, and they never see each other's. *Story (agency operator)* — "Keep my clients properly separated." **Where** — `/workspaces`, and the brand switcher at the top of the sidebar.

#### Steps

1. Sign in and complete onboarding once (§1.2) so one brand exists.
2. Open `/workspaces` → **Add Brand** → create a second, using a different test business.
3. Upload an asset to brand B (`/assets`).
4. Switch to brand A with the sidebar switcher and open `/assets`.

**Pass when** — brand A does not show brand B's asset. Calendar, settings and connections are likewise separate.

### 1.2 Onboarding — the five questions

**What it is** — five questions producing a brand "genome": what you can show, what you can film, what you are trying to achieve, whether anyone appears on camera. *Story (solo owner)* — "Set me up without a two-hour form." **Where** — `/onboarding`

#### Steps

1. Sign in with an account that has no brand — you are sent here automatically.
2. Answer all five using **Fernhill Roastery** from `TEST_BUSINESSES.md`.

3. Write down the answers you gave.
4. Finish and let it drop you into the app.

**Pass when** — you land in the app, not back at onboarding, and `/workspaces` lists the brand.

## 1.3 Build a brand from a website

---

**What it is** — point SPARK at a URL and it fills the brand in from the page. *Story (solo owner)* — "You can read what my business does." **Where** — onboarding, the "read my website" option.

### Steps

1. Start onboarding for a **new** brand.
2. Choose to supply a website instead of answering.
3. Paste a real small-business URL.
4. Read what comes back.

**Pass when** — name, category and one-liner are filled in and marked as *inferred* rather than confirmed, so you can correct them before continuing.

## 2. Brand setup — Settings → Brand Kits

---

Everything here applies to every post automatically.

### 2.1 Workspace logo

---

**What it is** — the logo used on rendered images and video. **Where** — Settings → Brand kits → *Workspace logo*

#### Steps

1. Press **Upload** and choose `docs/test-assets/fernhill-logo.png`.
2. Save.
3. Now clear it and press **Generate a placeholder** instead.

**Pass when** — the logo previews immediately after each action, and the generated one is a plain flat shape with no lettering.

### 2.2 Colour theme

---

**What it is** — three colours in order: background, text on it, accent. **Where** — Settings → Brand kits → *Color Theme*

#### Steps

1. Click a palette preset.
2. Now set the first two colours to something deliberately unreadable — e.g. mid grey text on light grey.
3. Look at the contrast strip underneath.

4. Fix it, and save.

**Pass when** — the strip shows your text colour on your background colour, so the unreadable pair is obvious on screen before it ever reaches a post.

## 2.3 Brand voice

---

**What it is** — four tone sliders plus the narration voice for video. *Story (creator)* — "It should sound like me."

**Where** — Settings → Brand kits → *Brand Voice*

### Steps

1. Set **formal 0.9 / playful 0.1** and save.
2. Generate a draft (§3.2) and keep the text.
3. Come back, set **formal 0.1 / playful 0.9**, save.
4. Generate another draft on the same playbook.
5. Compare the two.

**Pass when** — the writing is visibly different between them.

## 2.4 Typography

---

**What it is** — display and body typefaces for rendered posts. **Where** — Settings → Brand kits → *Typography Style*

### Steps

1. Pick a display face and a different body face.
2. Watch the sample block.
3. Save, then render an image post (§3.7).

**Pass when** — the sample renders in the chosen faces on your brand colours, and the rendered image uses them too.

## 2.5 Strict compliance

---

**What it is** — restricted topics, claims to avoid and banned words, plus a switch deciding whether a violation **blocks** a post or **holds** it. *Story (agency operator)* — "This client must never say that." **Where** — Settings → Brand kits → *Enable Strict Compliance*

### Steps

1. Add `guaranteed` to **Claims to avoid** and `artisanal` to banned words.
2. Turn strict mode **on**. Save.
3. Generate five drafts (§3.2).
4. Read each one for those words.

**Pass when** — neither word appears. If one ever does, the post is marked **blocked** with the reason on the item — it is not silently published.

## 2.6 Watermark

---

**What it is** — whether the logo is stamped on rendered media, and at what size and opacity. **Where** — Settings → Brand kits → *Watermark* (appears once a logo is set)

### Steps

1. With a logo set, confirm the Watermark card is visible.
2. Leave it on, render an image post (§3.7), look at the output.
3. Turn the watermark **off**, save, render again.

**Pass when** — the first render has the logo bottom-left, the second has none, and the logo is still set in §2.1 either way.

## 2.7 Brand knowledge

---

**What it is** — documents and pages SPARK may cite. Without them, guardrails hold any specific factual claim.

*Story (any)* — "Only say what's actually true about us." **Where** — Settings → Brand kits → *Brand Knowledge*

### Steps

1. Upload `docs/test-assets/fernhill-knowledge-base.pdf` (or `harbour-lane-knowledge-base.pdf`).
2. Wait for it to finish processing — it reports pages and chunks.
3. Generate a draft that would need a specific fact.
4. Now generate one that states a number the document does **not** contain.

**Pass when** — step 4 is blocked or flagged with a grounding reason, and the document counts toward the `knowledge` asset role so knowledge-dependent formats resolve.

**Pass when (grounding)** — step 3 cites a real number from the document rather than writing around it. A brand knowledge base holding "18 hours" and a "four-year-old starter" should produce copy containing those, not "a slow, delicious journey".

**Pass when (prohibitions)** — attaching the document also reports `restrictionsAdded`, and those rules then hold. A knowledge base with a "*things we must never claim*" section has its quoted phrases and claim categories added to **Enable Strict Compliance** (§2.5) automatically. Ask for a post that breaks one and it comes back held, naming the phrase.

The writer itself will still write the forbidden thing when asked directly — that cannot be prompted away. What stops it is the guardrail, which is why the document's rules are copied into governance rather than left as prose. Adding is one-way: nothing extracted can remove a rule you set by hand.

## 2.8 Offer

---

**What it is** — what you sell, and your primary call to action. **Where** — Settings → Brand kits → *Brand Knowledge* → Offer

### Steps

1. Set the primary CTA to `Start a subscription`. Save.
2. Generate a draft on a **product** playbook.
3. Generate one on an **educational** playbook.

**Pass when** — the product post carries the CTA; the educational one does not. Educational, proof and personality posts deliberately never pitch.

## 2.9 Templates

---

**What it is** — the layouts a post is built into: intros/outros, bumpers, caption presets, lower-thirds. **Where** — Settings → Brand kits → *Templates*

### Steps

1. Open each of the four tabs.
2. Add a template: give it a name and save it.
3. Open the **Versioning & approvals** tab.

**Pass when** — the template you added persists after a refresh, and the Versioning tab says **"Not built"** with an explanation. That is correct, not a defect.

## 2.10 Consent and avatar

---

**What it is** — whose face and voice SPARK may use, with a consent record, plus avatar and voice-clone configuration. *Story (creator)* — "Use my likeness, but only on the record." **Where** — Settings → Brand kits → *People and likeness*

### Steps

1. With **no** consent record, look at the available playbooks (§3.1).
2. Grant a consent record for a named person.
3. Look at the playbook list again.
4. Revoke the consent and look once more.

**Pass when** — avatar formats are unavailable at step 1, available at step 3, and unavailable again at step 4.

## 2.11 What this brand has learned

---

**What it is** — SPARK adapts the content mix from measured results; this shows what it has learned and lets you freeze it. **Where** — Settings → Brand kits → *What this brand has learned*

### Steps

1. Read the current confidence and what it is based on.
2. Press **Freeze**.
3. Refresh the page.

**Pass when** — it reports itself frozen and keeps what it already knew, rather than resetting to nothing.

---

## 3. Making content

### 3.1 Playbooks

---

**What it is** — around fifteen content formats; which you are offered is decided by your genome, never by your industry. *Story (creator)* — "Show me formats I can actually make." **Where** — Command Center ( /home ) → new post.

#### Steps

1. On **Fernhill Roastery** (physical craft), open the playbook list and note it.
2. Switch to **Northgate Actuarial** (data outcomes) and open the list again.
3. Open the **why** on any single playbook.

**Pass when** — the two lists differ meaningfully, and the why explains the choice in terms of what the brand can show or film — never "because you're a coffee shop".

### 3.2 Drafting a post

---

**What it is** — turns a chosen format into real copy, beat by beat, in your voice. **Where** — Command Center → new post → pick a playbook → draft.

#### Steps

1. Pick a playbook and press draft.
2. Optionally give an intent: `this week's Ethiopian arrival`.
3. Wait for every beat to fill.
4. Open the **Preview** tab and read the assembled post.

**Pass when** — every beat is filled at a sensible length, the post does not ask for the booking twice, and it ends with **hashtags specific to this post** — not `#smallbusiness` filler.

### 3.3 Regenerating one beat

---

**What it is** — rewrite just the hook, or just the CTA, leaving the rest alone.

#### Steps

1. In a drafted post, copy the text of two beats.

2. Regenerate **one** of them.
3. Compare both against what you copied.

**Pass when** — only the beat you asked for changed.

## 3.4 Variants and repurposing

---

**What it is** — several takes on one idea, or the same idea rebuilt as a different format for a different platform.

*Story (creator)* — "Turn one idea into a week of content."

### Steps

1. On a drafted post, generate **variants**.
2. Separately, **repurpose** a LinkedIn post into an X format.
3. Open the repurposed post's preview and count the hashtags.

**Pass when** — variants are genuinely different takes, and the repurposed X post carries **at most two** hashtags — its own, sized for X, not the LinkedIn original's.

## 3.5 Generated media

---

**What it is** — images, AI b-roll, avatar video, voiceover and dubbing, generated into individual beats. **Where** — Draft Panel → a beat → the generate actions.

### Steps

1. On an image playbook, generate an image into a beat.
2. On a video playbook, generate b-roll into one beat and a voiceover into another.
3. Note a beat's label and length, then dub that beat into another language.

**Pass when** — generated media replaces the beat and **keeps its timing and label**. The dub replaces the beat in place rather than adding a second one.

## 3.6 Storyboard editing

---

**What it is** — add, remove, retime and relabel scenes; add a lower-third; override the voice for one scene.

**Where** — Draft Panel → Storyboard.

### Steps

1. Note the running total at the top.
2. Insert a scene, then retime an existing one.
3. Add a lower-third name plate to one scene.
4. Now try to retime a scene so the video is far longer than the format allows.

**Pass when** — the running total updates each time, and step 4 is **refused with the reason**, not silently accepted.

## 3.7 Rendering

---

**What it is** — composes the beats into finished files, one per aspect ratio.

### Steps

1. With colours, type and logo set (§2.1–2.4), open a finished post.
2. Render it.
3. Open the output.

**Pass when** — output uses your brand colours (first colour as ground, second as the text on it), your typefaces, and your logo bottom-left.

---

## 4. Assets Library

### 4.1 Uploading and auto-captioning

---

**What it is** — your photos, video and documents, captioned automatically so they can be found by meaning.

**Where** — `/assets`

### Steps

1. Upload four or five images with meaningless filenames ( `IMG_4821.jpg` ).
2. Wait for captions to appear.
3. Search for something **in** one of the pictures, not its filename.

**Pass when** — search finds the image by its contents.

### 4.2 Folders and organisation

---

**What it is** — folders, moving, renaming, membership. *Story (agency operator)* — "Keep a client's material in order."

### Steps

1. Create a folder.
2. Move two assets into it.
3. Rename the folder, then refresh.

**Pass when** — the folder, its name and its contents all survive the refresh.

### 4.3 Roles and gaps

---

**What it is** — assets carry roles (social proof, product shot, work artifact...), and SPARK can say which missing role blocks the most formats.

### Steps

1. Set a role on one uploaded asset.
2. Open the gaps view for the brand.
3. Read the order of the list.

**Pass when** — gaps are ordered by how many formats each missing role would unlock, so the top of the list is what to shoot first.

## 4.4 Rights and reuse cooldown

---

**What it is** — a rights status per asset, and a cooldown so the same photo is not used every week.

### Steps

1. Set one asset's rights status to restricted.
2. Draft several posts in a row on the same playbook.
3. Watch which assets get picked.

**Pass when** — a restricted asset is not used, and an asset used in one draft is passed over for a fresh one in the next.

# 5. Campaigns and the calendar

---

## 5.1 Readiness check

---

**What it is** — before a campaign is created, a checklist of what would stop it producing anything, split into **blockers** and **warnings**. *Story (solo owner)* — "Tell me what's missing before I commit." **Where** — `/campaign/new`, on confirm.

### Steps

1. Use a brand with **no** connected account and **no** assets.
2. Start a campaign, choose an objective, and confirm.
3. Read the modal before doing anything else.

**Pass when** — the modal appears **before** the campaign is created. "Nowhere to publish" is a **blocker**; a missing asset is a **warning** carrying the number of extra posts it would unlock. The button reads **Create anyway** — you are not forced to fix anything.

## 5.2 Creating a campaign

---

**What it is** — an objective and a window; SPARK plans the mix and fills a calendar.

### Steps

1. Connect one account first (\$6).
2. Create a campaign with objective **sales** over 30 days.
3. Open `/calendar`.

4. Count how many posts are promotional versus educational.

**Pass when** — the calendar has posts on it, and they are spread across content pillars rather than being all promotional.

## 5.3 The calendar

---

**What it is** — the month, with drag-to-reschedule and slot recommendations. **Where** — `/calendar`

### Steps

1. Drag a post to a different day.
2. Refresh the page.
3. Now open the calendar of a brand with **no** campaign.

**Pass when** — the move survives the refresh, and the empty calendar explains *why* it is empty rather than just saying "nothing scheduled".

## 5.4 Publishing

---

**What it is** — posts go out on schedule, or immediately. **Where** — Draft Panel → Preview → Publish.

### Steps

1. Connect **Bluesky** (§6) — it needs only a handle and an app password.
2. Open a finished draft, press Publish, choose Bluesky.
3. Open the post on Bluesky itself.

**Pass when** — the post is live, the text matches the preview exactly (including hashtags), and the item in SPARK carries a receipt and a link back.

## 5.5 Rollback

---

**What it is** — take a published post back down.

### Steps

1. On the post from §5.4, press rollback and confirm.
2. Check Bluesky.
3. Look at the item in SPARK.

**Pass when** — gone from the platform, but the record is **kept and marked rolled back**, not deleted.

## 5.6 Tracked links

---

**What it is** — shortens your CTA link so clicks can be counted.

### Steps

1. On a draft, shorten the link.
2. Look at the Preview.
3. Publish and open the live post.

**Pass when** — the short link sits **after** the hashtags, and the platform still builds a preview card from it.

---

## 6. Connecting accounts

**What it is** — connect the platforms you publish to. **Where** — Settings → Account Connection

**Connectable today:** X, TikTok, LinkedIn, YouTube, Google Business (OAuth), and **Bluesky** (handle + app password, nothing to set up). **Not configured:** Instagram, Facebook, Threads, Pinterest, Reddit — `docs/VENDOR_SETUP.pdf` .

### Steps

1. Connect **Bluesky** first: your handle and an app password from your Bluesky settings. It needs no vendor setup, so it isolates "does connecting work".
2. Then connect an OAuth platform — X, TikTok or YouTube.
3. On the consent screen, note whether you are asked **which** account to use.
4. Back in SPARK, read the tile.
5. Leave it a while and re-open the page.

**Pass when** — step 3 asks which account; the tile shows **your** account name and picture rather than a generic platform name; health shows when the token expires; and the token refreshes on its own before it does.

---

## 7. Discovery

### 7.1 Trends

**What it is** — what is moving now, scored against your brand. *Story (creator)* — "What should I post about today?" **Where** — `/discovery`

### Steps

1. Open Discovery and let it load.
2. Open the **why** on the top-scoring trend.
3. Scan the list for anything wildly unrelated to the business.
4. Run `npx tsx --env-file-if-exists=apps/api/.env scripts/check-trend-sources.mts` .

**Pass when** — every trend carries a score and a why; nothing off-brand appears (it is removed, not shown faintly); and the script reports which sources are live. Four of eight are expected: YouTube, Hacker News, Product Hunt, Google Trends.

## 7.2 Turning a trend into a post

---

### Steps

1. Pick a trend with a decent score.
2. Press repurpose.
3. Read the resulting draft.

**Pass when** — the draft is about that trend, in your brand's voice, using a format your brand can actually make.

## 7.3 Watchlist, muting and influencers

---

### Steps

1. Add a topic to the watchlist.
2. Mute one source.
3. Reload Discovery.

**Pass when** — the watched topic is tracked and the muted source no longer contributes to the feed.

## 7.4 Hooks

---

**What it is** — opening lines drawn from what is working, adapted to your brand.

### Steps

1. Open the Hooks tab.
2. Take one hook into a draft.

**Pass when** — the hooks read as usable openers for *your* brand, not generic copywriting templates.

# 8. Automation recipes

**What it is** — standing rules that keep producing content: watch trends for keywords, pull an RSS feed, or import a CSV in bulk. *Story (agency operator)* — "Keep output going without me starting it." **Where** — `/automation`

### Steps

1. New recipe → **auto trend** → add two or three broad keywords → save.
2. Run it and open the output queue.
3. Now edit the recipe to use one absurd keyword ( `zzzqqq` ) and run again.
4. Approve one output and reject another.
5. Pause the recipe.

**Pass when** — step 2 produces outputs for review; step 3 returns a message saying *which* reason nothing came back — no trends matched those words, or they matched but scored too low for this brand; approve/reject each act on only that output; and a paused recipe produces nothing further.

---

## 9. Engagement Intelligence

**What it is** — incoming comments and DMs, classified, with drafted replies and sales opportunities pulled out. *Story (sales-driven brand)* — "Turn DMs into bookings." **Where** — `/agents?`

`tab=engagement`

### Steps

1. Open a conversation and read its classification.
2. Draft a reply and read it before doing anything else.
3. Check it has **not** been sent.
4. Approve it.
5. Open the sales opportunities list.

**Pass when** — the reply is in your brand's voice rather than generic; nothing is sent until step 4 unless auto-handling is explicitly on; opportunities are intent-scored and carry a recommended next action; anything risky is escalated rather than answered.

---

## 10. Capture — Direct + Finish

**What it is** — SPARK tells you exactly what to film, you send it over WhatsApp, and it finishes it into a post. *Story (solo owner)* — "I'll film it if you tell me what to point at."

### Steps

1. Pick a format that needs footage you do not have.
2. Start a capture session and read the brief.
3. Send a clip back over WhatsApp.
4. Now send a deliberately poor clip — very dark, or very short.

**Pass when** — the brief is specific and shootable (not "film something nice"); the sent media is ingested and attached to the post; and the poor clip **degrades to a simpler format** rather than failing outright.

---

## 11. Oversight and governance

### 11.1 Approval mode

---

**What it is** — autopublish, or hold everything for review. **Where** — Settings, and per campaign.

### Steps

1. Set approval mode to **review**.

2. Publish a post from a campaign.
3. Find it in the review queue and approve it.
4. Now create a post with **no campaign** and try to publish it.

**Pass when** — step 2 is held rather than published, step 3 publishes it, and step 4 is held **regardless of mode** — autonomy is a property of the campaign.

## 11.2 Approval rules

---

**What it is** — finer control: hold by platform, content type, risk or role.

### Steps

1. Add a rule holding everything for one specific platform.
2. Publish to that platform, then to another.

**Pass when** — only the named platform is held.

## 11.3 Review queue

---

### Steps

1. Create two or three held items.
2. Open the queue in the Command Center.
3. Act on one.

**Pass when** — everything waiting on a person is in one place, and acting on one removes only that item.

## 11.4 Guardrails

---

**What it is** — every draft is checked before scheduling and again before publishing: claim grounding, banned phrases, restricted topics, per-platform limits, duplicates, AI disclosure, rights, avatar saturation.

### Steps

1. With \$2.5 configured, generate a batch of drafts.
2. Find one that was blocked or flagged.
3. Open it.
4. Draft the same intent twice and look for the duplicate warning.

**Pass when** — a blocked post **stays visible with the reason on it**. The check never deletes the evidence.

## 11.5 Ask Spark and notifications

---

### Steps

1. Open Ask Spark and ask something about the current brand.

2. Trigger work that needs a decision from you.
3. Open the notification.

**Pass when** — the answer is grounded in this brand, and the work **waits** for your answer rather than guessing.

---

## 11.6 The agent

### Steps

1. Open agent status and pause it.
2. Wait past its normal interval.
3. Resume, then open the run log.

**Pass when** — nothing happens while paused, and each run in the log carries its reasoning.

---

## 12. Analytics

**What it is** — what happened after publishing: per-post metrics, brand trend over time, campaign reports, CTA click traffic, success metrics. **Where** — Command Center overview, and campaign reports.

### Steps

1. Publish at least one post (\$5.4).
2. Open the Command Center overview.
3. Open the campaign's report.

**Pass when** — the report compares what was **promised against what happened**, rather than only showing totals.

---

## 13. Team, roles and organisation

**What it is** — invite people, set roles, permissions and groups; set budgets and credit limits; configure SSO; read an audit log. *Story (agency operator)* — "Different people, different powers." **Where** — Settings → Team Roles, and Settings → Credit & Usage.

### Steps

1. Invite a second user as **viewer**.
2. As that user, try to publish a post.
3. Set a very low budget cap and generate several drafts.
4. Open the audit log.

**Pass when** — the viewer cannot publish; the budget **refuses spend** rather than overspending quietly; and every action appears in the audit log.

---

## 14. Agency features

**What it is** — a client roster, white-labelled links, and client proposals shareable on a public link.

**Where** — `/agency` , shared proposals at `/p/<token>` .

### Steps

1. Open the roster and check every brand is listed.
2. Draft a proposal for one client and share it.
3. Open the share link in a **private window**, signed out.
4. Accept or decline it there.

**Pass when** — the link opens with no login, shows the plan, and the decision is recorded back in SPARK.

---

## 15. Import, export and portability

**What it is** — export a brand whole, import it elsewhere. **Where** — Settings → Import / Export

### Steps

1. Export a fully set-up brand.
2. Create an empty brand.
3. Import the export into it.
4. Compare brand kit, knowledge and assets against the original.

**Pass when** — genome, rules, assets and knowledge all arrive, and the imported brand behaves like the original when you draft a post.

---

## Known limits

Correct behaviour today, not defects. Do not log these as failures.

| THING                                                            | WHY                                                         |
|------------------------------------------------------------------|-------------------------------------------------------------|
| Instagram, Facebook, Threads, Pinterest, Reddit will not connect | No app credentials set — <code>docs/VENDOR_SETUP.pdf</code> |
| LinkedIn publishing may be refused                               | Platform approval pending; the connection itself works      |
| 4 of 8 trend sources live                                        | Reddit, X, TikTok, Pinterest need keys                      |
| Brand kit <b>Versioning</b> tab says "Not built"                 | Nothing records a brand-kit version yet                     |
| Copy comes from the OpenAI fallback                              | The Anthropic organisation is disabled                      |
| Aggregator publishing off                                        | <code>AYRSHARE_API_KEY</code> unset; native adapters in use |

---

## Reporting a problem

Give the **section number**, what you did, what you saw, what you expected. For a publish failure the item itself carries the reason — copy that rather than "it didn't work". Browser console for screen problems; the API terminal for anything server-side.