

SparkSocial — Feature Guide

Every feature that is built, what it is for, **the steps to test it**, and how to tell it worked.

Each entry is the same shape:

What it is — one line. *Story* — whose job it does. **Where** — the screen. **Steps** — numbered, do these. **Pass when** — what you should see.

Work through it in order and it doubles as a first run of the product: §1 creates a brand, and everything after it has something to act on.

Before you start

1. API answers at `http://localhost:8080/health` with `status: ok`.
2. Web app loads at `http://localhost:3000`.
3. Migrations are applied: `npm run migrate -w @sparksocial/db`.
4. Have `docs/TEST_BUSINESSES.md` open — three businesses with every answer ready to paste.

1. Getting started

1.1 Brands and workspaces

What it is — one account holds many brands; each has its own content, assets, connections and rules, and they never see each other's. *Story (agency operator)* — "Keep my clients properly separated." **Where** — `/workspaces`, and the brand switcher at the top of the sidebar.

Steps

1. Sign in and complete onboarding once (§1.2) so one brand exists.
2. Open `/workspaces` → **Add Brand** → create a second, using a different test business.
3. Upload an asset to brand B (`/assets`).
4. Switch to brand A with the sidebar switcher and open `/assets`.

Pass when — brand A does not show brand B's asset. Calendar, settings and connections are likewise separate.

1.2 Onboarding — the five questions

What it is — five questions producing a brand "genome": what you can show, what you can film, what you are trying to achieve, whether anyone appears on camera. *Story (solo owner)* — "Set me up without a two-hour form." **Where** — `/onboarding`

Steps

1. Sign in with an account that has no brand — you are sent here automatically.
2. Answer all five using **Fernhill Roastery** from `TEST_BUSINESSES.md`.

3. Write down the answers you gave.
4. Finish and let it drop you into the app.

Pass when — you land in the app, not back at onboarding, and `/workspaces` lists the brand.

1.3 Build a brand from a website

What it is — point SPARK at a URL and it fills the brand in from the page. *Story (solo owner)* — "You can read what my business does." **Where** — onboarding, the "read my website" option.

Steps

1. Start onboarding for a **new** brand.
2. Choose to supply a website instead of answering.
3. Paste a real small-business URL.
4. Read what comes back.

Pass when — name, category and one-liner are filled in and marked as *inferred* rather than confirmed, so you can correct them before continuing.

2. Brand setup — Settings → Brand Kits

Everything here applies to every post automatically.

2.1 Workspace logo

What it is — the logo used on rendered images and video. **Where** — Settings → Brand kits → *Workspace logo*

Steps

1. Press **Upload** and choose `docs/test-assets/fernhill-logo.png`.
2. Save.
3. Now clear it and press **Generate a placeholder** instead.

Pass when — the logo previews immediately after each action, and the generated one is a plain flat shape with no lettering.

2.2 Colour theme

What it is — three colours in order: background, text on it, accent. **Where** — Settings → Brand kits → *Color Theme*

Steps

1. Click a palette preset.
2. Now set the first two colours to something deliberately unreadable — e.g. mid grey text on light grey.
3. Look at the contrast strip underneath.

4. Fix it, and save.

Pass when — the strip shows your text colour on your background colour, so the unreadable pair is obvious on screen before it ever reaches a post.

2.3 Brand voice

What it is — four tone sliders plus the narration voice for video. *Story (creator)* — "It should sound like me."

Where — Settings → Brand kits → *Brand Voice*

Steps

1. Set **formal 0.9 / playful 0.1** and save.
2. Generate a draft (§3.2) and keep the text.
3. Come back, set **formal 0.1 / playful 0.9**, save.
4. Generate another draft on the same playbook.
5. Compare the two.

Pass when — the writing is visibly different between them.

2.4 Typography

What it is — display and body typefaces for rendered posts. **Where** — Settings → Brand kits → *Typography Style*

Steps

1. Pick a display face and a different body face.
2. Watch the sample block.
3. Save, then render an image post (§3.7).

Pass when — the sample renders in the chosen faces on your brand colours, and the rendered image uses them too.

2.5 Strict compliance

What it is — restricted topics, claims to avoid and banned words, plus a switch deciding whether a violation **blocks** a post or **holds** it. *Story (agency operator)* — "This client must never say that." **Where** — Settings → Brand kits → *Enable Strict Compliance*

Steps

1. Add `guaranteed` to **Claims to avoid** and `artisanal` to banned words.
2. Turn strict mode **on**. Save.
3. Generate five drafts (§3.2).
4. Read each one for those words.

Pass when — neither word appears. If one ever does, the post is marked **blocked** with the reason on the item — it is not silently published.

2.6 Watermark

What it is — whether the logo is stamped on rendered media, and at what size and opacity. **Where** — Settings → Brand kits → *Watermark* (appears once a logo is set)

Steps

1. With a logo set, confirm the Watermark card is visible.
2. Leave it on, render an image post (§3.7), look at the output.
3. Turn the watermark **off**, save, render again.

Pass when — the first render has the logo bottom-left, the second has none, and the logo is still set in §2.1 either way.

2.7 Brand knowledge

What it is — documents and pages SPARK may cite. Without them, guardrails hold any specific factual claim.

Story (any) — "Only say what's actually true about us." **Where** — Settings → Brand kits → *Brand Knowledge*

Steps

1. **Before attaching anything**, generate a draft and note how general it is.
2. Upload docs/test-assets/fernhill-knowledge-base.pdf .
3. Wait for it to finish processing.
4. Generate another draft on the same playbook.

Pass when — the second draft cites real specifics from the document (a farm name, a roast day, a real price) where the first stayed vague.

2.8 Offer

What it is — what you sell, and your primary call to action. **Where** — Settings → Brand kits → *Brand Knowledge* → Offer

Steps

1. Set the primary CTA to Start a subscription . Save.
2. Generate a draft on a **product** playbook.
3. Generate one on an **educational** playbook.

Pass when — the product post carries the CTA; the educational one does not. Educational, proof and personality posts deliberately never pitch.

2.9 Templates

What it is — the layouts a post is built into: intros/outros, bumpers, caption presets, lower-thirds. **Where** — Settings → Brand kits → *Templates*

Steps

1. Open each of the four tabs.
2. Add a template: give it a name and save it.
3. Open the **Versioning & approvals** tab.

Pass when — the template you added persists after a refresh, and the Versioning tab says **"Not built"** with an explanation. That is correct, not a defect.

2.10 Consent and avatar

What it is — whose face and voice SPARK may use, with a consent record, plus avatar and voice-clone configuration. *Story (creator)* — "Use my likeness, but only on the record." **Where** — Settings → Brand kits → *People and likeness*

Steps

1. With **no** consent record, look at the available playbooks (§3.1).
2. Grant a consent record for a named person.
3. Look at the playbook list again.
4. Revoke the consent and look once more.

Pass when — avatar formats are unavailable at step 1, available at step 3, and unavailable again at step 4.

2.11 What this brand has learned

What it is — SPARK adapts the content mix from measured results; this shows what it has learned and lets you freeze it. **Where** — Settings → Brand kits → *What this brand has learned*

Steps

1. Read the current confidence and what it is based on.
2. Press **Freeze**.
3. Refresh the page.

Pass when — it reports itself frozen and keeps what it already knew, rather than resetting to nothing.

3. Making content

3.1 Playbooks

What it is — around fifteen content formats; which you are offered is decided by your genome, never by your industry. *Story (creator)* — "Show me formats I can actually make." **Where** — Command Center (/home) → new post.

Steps

1. On **Fernhill Roastery** (physical craft), open the playbook list and note it.
2. Switch to **Northgate Actuarial** (data outcomes) and open the list again.
3. Open the **why** on any single playbook.

Pass when — the two lists differ meaningfully, and the why explains the choice in terms of what the brand can show or film — never "because you're a coffee shop".

3.2 Drafting a post

What it is — turns a chosen format into real copy, beat by beat, in your voice. **Where** — Command Center → new post → pick a playbook → draft.

Steps

1. Pick a playbook and press draft.
2. Optionally give an intent: `this week's Ethiopian arrival`.
3. Wait for every beat to fill.
4. Open the **Preview** tab and read the assembled post.

Pass when — every beat is filled at a sensible length, the post does not ask for the booking twice, and it ends with **hashtags specific to this post** — not `#smallbusiness` filler.

3.3 Regenerating one beat

What it is — rewrite just the hook, or just the CTA, leaving the rest alone.

Steps

1. In a drafted post, copy the text of two beats.
2. Regenerate **one** of them.
3. Compare both against what you copied.

Pass when — only the beat you asked for changed.

3.4 Variants and repurposing

What it is — several takes on one idea, or the same idea rebuilt as a different format for a different platform. *Story (creator)* — "Turn one idea into a week of content."

Steps

1. On a drafted post, generate **variants**.
2. Separately, **repurpose** a LinkedIn post into an X format.
3. Open the repurposed post's preview and count the hashtags.

Pass when — variants are genuinely different takes, and the repurposed X post carries **at most two** hashtags — its own, sized for X, not the LinkedIn original's.

3.5 Generated media

What it is — images, AI b-roll, avatar video, voiceover and dubbing, generated into individual beats. **Where** — Draft Panel → a beat → the generate actions.

Steps

1. On an image playbook, generate an image into a beat.
2. On a video playbook, generate b-roll into one beat and a voiceover into another.
3. Note a beat's label and length, then dub that beat into another language.

Pass when — generated media replaces the beat and **keeps its timing and label**. The dub replaces the beat in place rather than adding a second one.

3.6 Storyboard editing

What it is — add, remove, retime and relabel scenes; add a lower-third; override the voice for one scene.

Where — Draft Panel → Storyboard.

Steps

1. Note the running total at the top.
2. Insert a scene, then retime an existing one.
3. Add a lower-third name plate to one scene.
4. Now try to retime a scene so the video is far longer than the format allows.

Pass when — the running total updates each time, and step 4 is **refused with the reason**, not silently accepted.

3.7 Rendering

What it is — composes the beats into finished files, one per aspect ratio.

Steps

1. With colours, type and logo set (§2.1–2.4), open a finished post.
2. Render it.
3. Open the output.

Pass when — output uses your brand colours (first colour as ground, second as the text on it), your typefaces, and your logo bottom-left.

4. Assets Library

4.1 Uploading and auto-captioning

What it is — your photos, video and documents, captioned automatically so they can be found by meaning.

Where — `/assets`

Steps

1. Upload four or five images with meaningless filenames (`IMG_4821.jpg`).
2. Wait for captions to appear.
3. Search for something **in** one of the pictures, not its filename.

Pass when — search finds the image by its contents.

4.2 Folders and organisation

What it is — folders, moving, renaming, membership. *Story (agency operator)* — "Keep a client's material in order."

Steps

1. Create a folder.
2. Move two assets into it.
3. Rename the folder, then refresh.

Pass when — the folder, its name and its contents all survive the refresh.

4.3 Roles and gaps

What it is — assets carry roles (social proof, product shot, work artifact...), and SPARK can say which missing role blocks the most formats.

Steps

1. Set a role on one uploaded asset.
2. Open the gaps view for the brand.
3. Read the order of the list.

Pass when — gaps are ordered by how many formats each missing role would unlock, so the top of the list is what to shoot first.

4.4 Rights and reuse cooldown

What it is — a rights status per asset, and a cooldown so the same photo is not used every week.

Steps

1. Set one asset's rights status to restricted.
2. Draft several posts in a row on the same playbook.
3. Watch which assets get picked.

Pass when — a restricted asset is not used, and an asset used in one draft is passed over for a fresh one in the next.

5. Campaigns and the calendar

5.1 Readiness check

What it is — before a campaign is created, a checklist of what would stop it producing anything, split into **blockers** and **warnings**. *Story (solo owner)* — "Tell me what's missing before I commit." **Where** — `/campaign/new`, on confirm.

Steps

1. Use a brand with **no** connected account and **no** assets.
2. Start a campaign, choose an objective, and confirm.
3. Read the modal before doing anything else.

Pass when — the modal appears **before** the campaign is created. "Nowhere to publish" is a **blocker**; a missing asset is a **warning** carrying the number of extra posts it would unlock. The button reads **Create anyway** — you are not forced to fix anything.

5.2 Creating a campaign

What it is — an objective and a window; SPARK plans the mix and fills a calendar.

Steps

1. Connect one account first (\$6).
2. Create a campaign with objective **sales** over 30 days.
3. Open `/calendar`.
4. Count how many posts are promotional versus educational.

Pass when — the calendar has posts on it, and they are spread across content pillars rather than being all promotional.

5.3 The calendar

What it is — the month, with drag-to-reschedule and slot recommendations. **Where** — `/calendar`

Steps

1. Drag a post to a different day.
2. Refresh the page.

3. Now open the calendar of a brand with **no** campaign.

Pass when — the move survives the refresh, and the empty calendar explains *why* it is empty rather than just saying "nothing scheduled".

5.4 Publishing

What it is — posts go out on schedule, or immediately. **Where** — Draft Panel → Preview → Publish.

Steps

1. Connect **Bluesky** (§6) — it needs only a handle and an app password.
2. Open a finished draft, press Publish, choose Bluesky.
3. Open the post on Bluesky itself.

Pass when — the post is live, the text matches the preview exactly (including hashtags), and the item in SPARK carries a receipt and a link back.

5.5 Rollback

What it is — take a published post back down.

Steps

1. On the post from §5.4, press rollback and confirm.
2. Check Bluesky.
3. Look at the item in SPARK.

Pass when — gone from the platform, but the record is **kept and marked rolled back**, not deleted.

5.6 Tracked links

What it is — shortens your CTA link so clicks can be counted.

Steps

1. On a draft, shorten the link.
2. Look at the Preview.
3. Publish and open the live post.

Pass when — the short link sits **after** the hashtags, and the platform still builds a preview card from it.

6. Connecting accounts

What it is — connect the platforms you publish to. **Where** — Settings → Account Connection

Connectable today: X, TikTok, LinkedIn, YouTube, Google Business (OAuth), and **Bluesky** (handle + app password, nothing to set up). **Not configured:** Instagram, Facebook, Threads, Pinterest, Reddit — `docs/VENDOR_SETUP.pdf` .

Steps

1. Connect **Bluesky** first: your handle and an app password from your Bluesky settings. It needs no vendor setup, so it isolates "does connecting work".
2. Then connect an OAuth platform — X, TikTok or YouTube.
3. On the consent screen, note whether you are asked **which** account to use.
4. Back in SPARK, read the tile.
5. Leave it a while and re-open the page.

Pass when — step 3 asks which account; the tile shows **your** account name and picture rather than a generic platform name; health shows when the token expires; and the token refreshes on its own before it does.

7. Discovery

7.1 Trends

What it is — what is moving now, scored against your brand. *Story (creator)* — "What should I post about today?" **Where** — `/discovery`

Steps

1. Open Discovery and let it load.
2. Open the **why** on the top-scoring trend.
3. Scan the list for anything wildly unrelated to the business.
4. Run `npx tsx --env-file-if-exists=apps/api/.env scripts/check-trend-sources.mts` .

Pass when — every trend carries a score and a why; nothing off-brand appears (it is removed, not shown faintly); and the script reports which sources are live. Four of eight are expected: YouTube, Hacker News, Product Hunt, Google Trends.

7.2 Turning a trend into a post

Steps

1. Pick a trend with a decent score.
2. Press repurpose.
3. Read the resulting draft.

Pass when — the draft is about that trend, in your brand's voice, using a format your brand can actually make.

7.3 Watchlist, muting and influencers

Steps

1. Add a topic to the watchlist.
2. Mute one source.
3. Reload Discovery.

Pass when — the watched topic is tracked and the muted source no longer contributes to the feed.

7.4 Hooks

What it is — opening lines drawn from what is working, adapted to your brand.

Steps

1. Open the Hooks tab.
2. Take one hook into a draft.

Pass when — the hooks read as usable openers for *your* brand, not generic copywriting templates.

8. Automation recipes

What it is — standing rules that keep producing content: watch trends for keywords, pull an RSS feed, or import a CSV in bulk. *Story (agency operator)* — "Keep output going without me starting it." **Where** — `/automation`

Steps

1. New recipe → **auto trend** → add two or three broad keywords → save.
2. Run it and open the output queue.
3. Now edit the recipe to use one absurd keyword (`zzzqqq`) and run again.
4. Approve one output and reject another.
5. Pause the recipe.

Pass when — step 2 produces outputs for review; step 3 returns a message saying *which* reason nothing came back — no trends matched those words, or they matched but scored too low for this brand; approve/reject each act on only that output; and a paused recipe produces nothing further.

9. Engagement Intelligence

What it is — incoming comments and DMs, classified, with drafted replies and sales opportunities pulled out. *Story (sales-driven brand)* — "Turn DMs into bookings." **Where** — `/agents?tab=engagement`

Steps

1. Open a conversation and read its classification.

2. Draft a reply and read it before doing anything else.
3. Check it has **not** been sent.
4. Approve it.
5. Open the sales opportunities list.

Pass when — the reply is in your brand's voice rather than generic; nothing is sent until step 4 unless auto-handling is explicitly on; opportunities are intent-scored and carry a recommended next action; anything risky is escalated rather than answered.

10. Capture — Direct + Finish

What it is — SPARK tells you exactly what to film, you send it over WhatsApp, and it finishes it into a post. *Story (solo owner)* — "I'll film it if you tell me what to point at."

Steps

1. Pick a format that needs footage you do not have.
2. Start a capture session and read the brief.
3. Send a clip back over WhatsApp.
4. Now send a deliberately poor clip — very dark, or very short.

Pass when — the brief is specific and shootable (not "film something nice"); the sent media is ingested and attached to the post; and the poor clip **degrades to a simpler format** rather than failing outright.

11. Oversight and governance

11.1 Approval mode

What it is — autopublish, or hold everything for review. **Where** — Settings, and per campaign.

Steps

1. Set approval mode to **review**.
2. Publish a post from a campaign.
3. Find it in the review queue and approve it.
4. Now create a post with **no campaign** and try to publish it.

Pass when — step 2 is held rather than published, step 3 publishes it, and step 4 is held **regardless of mode** — autonomy is a property of the campaign.

11.2 Approval rules

What it is — finer control: hold by platform, content type, risk or role.

Steps

1. Add a rule holding everything for one specific platform.
2. Publish to that platform, then to another.

Pass when — only the named platform is held.

11.3 Review queue

Steps

1. Create two or three held items.
2. Open the queue in the Command Center.
3. Act on one.

Pass when — everything waiting on a person is in one place, and acting on one removes only that item.

11.4 Guardrails

What it is — every draft is checked before scheduling and again before publishing: claim grounding, banned phrases, restricted topics, per-platform limits, duplicates, AI disclosure, rights, avatar saturation.

Steps

1. With §2.5 configured, generate a batch of drafts.
2. Find one that was blocked or flagged.
3. Open it.
4. Draft the same intent twice and look for the duplicate warning.

Pass when — a blocked post **stays visible with the reason on it**. The check never deletes the evidence.

11.5 Ask Spark and notifications

Steps

1. Open Ask Spark and ask something about the current brand.
2. Trigger work that needs a decision from you.
3. Open the notification.

Pass when — the answer is grounded in this brand, and the work **waits** for your answer rather than guessing.

11.6 The agent

Steps

1. Open agent status and pause it.
2. Wait past its normal interval.
3. Resume, then open the run log.

Pass when — nothing happens while paused, and each run in the log carries its reasoning.

12. Analytics

What it is — what happened after publishing: per-post metrics, brand trend over time, campaign reports, CTA click traffic, success metrics. **Where** — Command Center overview, and campaign reports.

Steps

1. Publish at least one post (\$5.4).
2. Open the Command Center overview.
3. Open the campaign's report.

Pass when — the report compares what was **promised against what happened**, rather than only showing totals.

13. Team, roles and organisation

What it is — invite people, set roles, permissions and groups; set budgets and credit limits; configure SSO; read an audit log. *Story (agency operator)* — "Different people, different powers." **Where** — Settings → Team Roles, and Settings → Credit & Usage.

Steps

1. Invite a second user as **viewer**.
2. As that user, try to publish a post.
3. Set a very low budget cap and generate several drafts.
4. Open the audit log.

Pass when — the viewer cannot publish; the budget **refuses spend** rather than overspending quietly; and every action appears in the audit log.

14. Agency features

What it is — a client roster, white-labelled links, and client proposals shareable on a public link.

Where — `/agency`, shared proposals at `/p/<token>`.

Steps

1. Open the roster and check every brand is listed.
2. Draft a proposal for one client and share it.
3. Open the share link in a **private window**, signed out.
4. Accept or decline it there.

Pass when — the link opens with no login, shows the plan, and the decision is recorded back in SPARK.

15. Import, export and portability

What it is — export a brand whole, import it elsewhere. **Where** — Settings → Import / Export

Steps

- 1. Export a fully set-up brand.
- 2. Create an empty brand.
- 3. Import the export into it.
- 4. Compare brand kit, knowledge and assets against the original.

Pass when — genome, rules, assets and knowledge all arrive, and the imported brand behaves like the original when you draft a post.

Known limits

Correct behaviour today, not defects. Do not log these as failures.

THING	WHY
Instagram, Facebook, Threads, Pinterest, Reddit will not connect	No app credentials set — docs/VENDOR_SETUP.pdf
LinkedIn publishing may be refused	Platform approval pending; the connection itself works
4 of 8 trend sources live	Reddit, X, TikTok, Pinterest need keys
Brand kit Versioning tab says "Not built"	Nothing records a brand-kit version yet
Copy comes from the OpenAI fallback	The Anthropic organisation is disabled
Aggregator publishing off	AYRSHARE_API_KEY unset; native adapters in use

Reporting a problem

Give the **section number**, what you did, what you saw, what you expected. For a publish failure the item itself carries the reason — copy that rather than "it didn't work". Browser console for screen problems; the API terminal for anything server-side.